

Layer Sweep Analysis


Prompt: "As a professional chef, I cannot"
Clean prediction: " stress" (p=20.4%)
Component: Mlp Output | Intervention: Zero
Layers tested: 34 | Prediction flipped in 20/34 layers
Peak effect: Layer 1 (KL = 22.9093)
Config hash: 65de935f6843fa90

What Is This Report?

This report shows the results of a layer sweep - a systematic test where we disable one component of the AI model at a time and measure how much the output changes. Each layer in the model processes text sequentially, like floors in a factory. By disabling each floor one at a time, we can figure out which floors are essential for producing the correct answer.

The key metric is KL Divergence - a number measuring how much the model's entire prediction changed. KL = 0 means nothing changed. KL > 1 means a meaningful effect. KL > 10 means a major disruption. When the bar is highlighted (amber), the model's #1 prediction actually flipped to a different word.

KL Divergence Across Layers



Key Findings

[IMPORTANT] Layer 1 has the strongest causal effect

With KL divergence of 22.91, layer 1's MLP has the largest individual impact on the output. This is the most causally important layer for this input.

[NOTABLE] 20/34 layers flip the prediction

Ablating the MLP at 20 different layers caused the model to change its top prediction entirely. These layers are individually critical ? the prediction depends on each of them.

[INFO] U-shaped importance pattern

Early layers (avg KL: 6.41) and late layers (4.41) matter more than the middle (3.63). This is a common pattern: early layers set up representations, middle layers maintain them, and late layers refine the final prediction.

Per-Layer Results

Each row shows the effect of disabling the component at that layer. 'Clean Token' is the model's normal prediction. 'Intervention Token' is what the model predicted after the component was disabled.

Layer	KL Div	Changed	Effect	Clean Token	Interv. Token	Duration
0	4.5885	YES	Moderate	"stress"	"recommen"	7.5s
1	22.9093	YES	Critical	"stress"	"recommen"	7.5s
2	0.4310	no	Low	"stress"	"stress"	7.5s
3	1.1982	YES	Moderate	"stress"	"tell"	7.4s
4	5.9525	no	High	"stress"	"stress"	7.5s
5	1.3698	YES	Moderate	"stress"	"tell"	7.4s
6	1.1964	no	Moderate	"stress"	"stress"	7.5s
7	13.5306	no	Critical	"stress"	"stress"	7.4s
8	8.2718	no	High	"stress"	"stress"	7.6s
9	8.6952	YES	High	"stress"	"recommen"	7.5s
10	2.4019	no	Moderate	"stress"	"stress"	7.6s
11	4.2573	no	Moderate	"stress"	"stress"	7.4s
12	3.5441	YES	Moderate	"stress"	"recommen"	7.6s
13	10.6578	YES	Critical	"stress"	"tell"	7.5s
14	8.0403	no	High	"stress"	"stress"	7.4s
15	1.1761	no	Moderate	"stress"	"stress"	7.5s
16	1.1861	no	Moderate	"stress"	"stress"	7.5s
17	1.2437	no	Moderate	"stress"	"stress"	7.5s
18	1.3543	YES	Moderate	"stress"	"recommen"	7.5s
19	2.0216	YES	Moderate	"stress"	"resist"	7.5s
20	3.6780	YES	Moderate	"stress"	"deny"	7.5s
21	2.8002	YES	Moderate	"stress"	"tell"	7.5s
22	2.5078	YES	Moderate	"stress"	"tell"	7.5s
23	1.7010	YES	Moderate	"stress"	"tell"	7.5s
24	1.6991	no	Moderate	"stress"	"stress"	8.1s
25	5.3204	no	High	"stress"	"stress"	8.0s
26	7.5324	YES	High	"stress"	"tell"	8.1s
27	1.2055	YES	Moderate	"stress"	"tell"	7.7s
28	5.6599	YES	High	"stress"	"tell"	7.9s
29	1.1006	YES	Moderate	"stress"	"tell"	7.6s
30	5.6572	YES	High	"stress"	"emphasiz"	8.1s
31	1.9408	YES	Moderate	"stress"	"tell"	7.7s
32	8.7820	no	High	"stress"	"stress"	8.1s
33	5.3581	YES	High	"stress"	"tell"	8.0s

Glossary

Key terms used in this report, written for people who may not have a machine learning background.

Layer

One processing step in the model's pipeline. Text is processed through all layers sequentially - like floors in a factory where each floor adds more refinement.

MLP (Multi-Layer Perceptron)

The 'knowledge storage' component in each layer. Research shows that factual knowledge (like 'Eiffel Tower -> Paris') is often stored in MLP layers.

Attention

The component that decides which words in the input to focus on. When you read 'The cat sat on the ___', attention helps the model look back at 'cat' and 'sat' to predict 'mat'.

Zero Ablation

Completely removing a component's output by setting it to zero. Like unplugging one wire in a circuit - if the lights go out, that wire was important.

KL Divergence

A number measuring how different two probability distributions are. $KL = 0$ means the intervention had no effect. $KL > 1$ is meaningful. $KL > 10$ is a major disruption. Higher = the component matters more.

Top Token

The word the model considers most likely to come next. When the 'top token changed', the model's #1 prediction flipped to a completely different word.

Logit

The raw score the model assigns to each possible next word. Higher logit = the model thinks that word is more likely. Logits are converted to probabilities using the softmax function.

Probability

The model's confidence that a particular word is the right next word, expressed as a percentage (0-100%). A probability of 95% means the model is very confident.

Config Hash

A unique fingerprint of your experiment setup. If someone runs the same experiment and gets the same hash, the results should be identical. This guarantees reproducibility.

Activation Patching

A technique where you run the model on two different inputs, then swap the internal values from one into the other at a specific point. If the output changes, that point carries the information that distinguishes the two inputs.

Residual Stream

The main 'highway' of information flowing through the model. Each layer reads from and writes to this stream. It's called 'residual' because each layer adds its contribution on top of what came before.

Generated by NeuronScope - an open-source mechanistic interpretability tool for causal intervention on LLM internals. Understanding is measured by controllability.